

Mining Big Data - Final Report

Investigating Paragraph Vectors

a1632538 Zachary Forman

a1646930 James Caddy

October 30, 2016

Introduction

We live in a world where enormous amounts of textual data exists, and even more is generated each day. This data presents both a major opportunity and major challenges, being plentiful and rich in information, but also high dimensional and loosely structured. We can view text as consisting of words, and can describe any collection of words as a document. In the past, textual documents are frequently viewed as just a collection of words, the so-called “Bag of Words” model [1]. Each word in such a model is typically just a number - the semantic connotations of words and documents discarded for the sake of efficiency.

In 2013, Mikolov *et al.* presented the **word2vec** algorithm [2]. This approach allowed for each word to be represented by a vector that contains a compact representation of the semantic meaning of that word. In 2014, Le and Mikolov extended this approach [3] with the **doc2vec** algorithm to allow learning vectors that summarize the semantic meaning of a document.

We seek to extend Le and Mikolov’s work, reproducing their **doc2vec** results and further extending them to the domains of document similarity and understanding source code.

Contributions

This work makes the following contributions:

- Replication of a relatively recent paper [3] that makes some contentious claims [4].
- Comparison of **doc2vec**’s suitability as a document similarity measure as compared to Jaccard similarity.
- Application of the **doc2vec** algorithm to a novel domain – source code.

Background

Word Vectors

The first meaningful attempt at learning vector representations was Rumelhart *et al.* ’s 1986 contribution [5]. This work, while relatively primitive, paved the way for more sophisticated models, such as Bengio *et al.* ’s 2003 work [6]. This work presents a model that is extremely similar to the current state of the art: Learn vectors for words based on errors seen while trying to predict surrounding context.

Mikolov *et al.* present various enhancements to this technique [7, 2], improving both the speed of training the word vectors and the quality of the learned representations. The algorithm presented in these papers is known as **word2vec**, and works as follows (see also Figure 1):

1. First, a set of sequential words is taken, and the central word set aside.
2. This text is then projected into the vector space, and summed.
3. The resulting vector is then used to predict the central word.
4. The error in the prediction is backpropagated to refine both the estimates of the word vectors and the prediction network.

Mikolov *et al.* show that these learnt word vectors not only retain semantic information, but also relational information; for example, the result of $\text{vec}(\text{king}) - \text{vec}(\text{man}) + \text{vec}(\text{woman}) \approx \text{vec}(\text{queen})$.

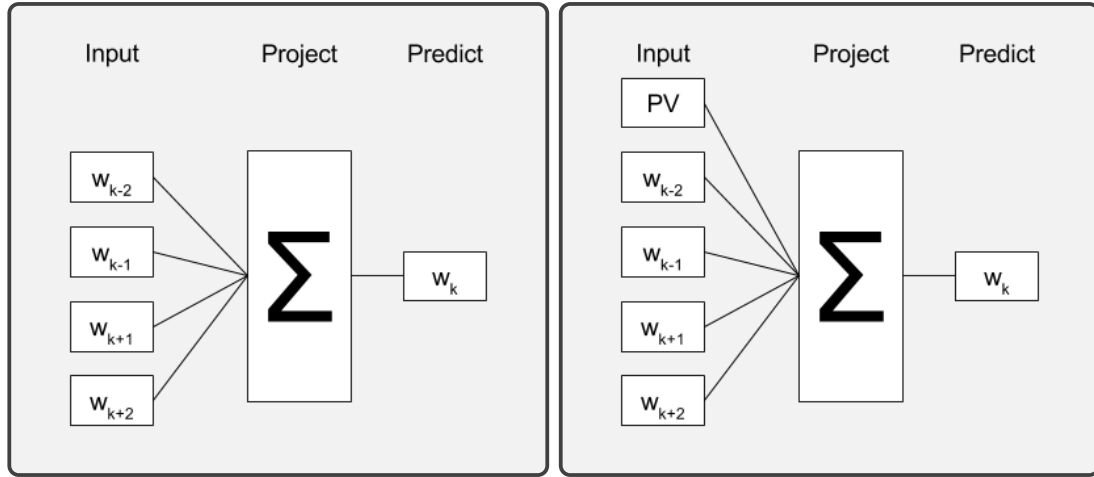


Figure 1: Graphical representation of the `word2vec` (left) and `doc2vec` (right) algorithms.

Paragraph Vectors

In 2014, Le and Mikolov presented a method of altering the `word2vec` algorithm to jointly train vector representations of documents in addition to word vectors [3]. This method is based on the idea of basically adding a “context vector”, representing an individual document, to each optimization. The process of learning these document vectors is equivalent to that of learning word vectors in the `word2vec` algorithm, except an additional paragraph vector (representing the document a given text sample comes from) is added - see Figure 1 for a graphical representation of the difference. This modified `word2vec` algorithm is called `doc2vec`.

Further exploration by Lau *et al.* shows that the `doc2vec` algorithm provides many of the same interesting properties that `word2vec` can provide [8]. Not only do the document vectors capture semantic information, they *also* capture relations: famous singers are located in similar places, and using document vectors trained on wikipedia pages, “Lady Gaga” - “America” + “Japan” is closest to “Ayumi Hamasaki”, a famous Japanese singer [9].

Data Collection

Data Sources

We used three sources of data. Firstly, we use Maas’ IMDB set [10] for verification and English language documents which can be found at <http://ai.stanford.edu/~amaas/data/sentiment/>. Secondly, we used The Wikimedia Foundation’s English Wikipedia data dump [11] as a source for English language documents. We downloaded the full wiki dump from <https://dumps.wikimedia.org/enwiki/20160820/>. Lastly, we used The Chromium Project’s Chromium repository [12] as a source for source code documents. We cloned the repository at revision 4fe31bb06cf458234d7017950a8b2b82427487c8.

Data Scale

The IMDB dataset contains a total of 100,000 files. Half are labelled as either a positive or negative reviews, and the remainder are unlabelled. It is 346MB in size and contains 23 million words, from a vocabulary of 90 thousand words.

The English Wikipedia Dataset is a single XML file, 13GB compressed and 55GB uncompressed. It contains 16 million pages, however this includes redirects and meta-pages such as categories and portals, and therefore only 5 million of these are useful as documents. Excluding markup, this dataset contains approximately 4 billion words.

The Chromium codebase is 2.6GB in size, and contains 245 thousand files. Of these, 23 thousand are C and C++ files, which contain 4.8 million lines of code and 608 thousand lines of comments.

Data Format

The fundamental requirements for the `doc2vec` [3] algorithm to work are that each document can be associated with the words in the document, and that words in each document are just a series of words, i.e. no excess punctuation or markup.

This requirement can easily be met by storing our set of documents as a newline separated list of space separated tokens. The following example shows how a simple document set might be transformed into this file format:

```
{ "My cow was sick, but now it is better.", "Apples are nice.", "Amazing results!" }
```

Convert all text to lowercase and remove all punctuation:

```
{ "my cow was sick but now it is better", "apples are nice", "amazing results" }
```

Concatenate documents, separated by newlines:

```
my cow was sick but now it is better
apples are nice
amazing results
```

This format is simple to produce, simple to parse, and - given that the `doc2vec` algorithm needs to process every word anyway - is efficient to read. Perhaps more importantly, especially for large datasets like the Chromium dataset, this storage format has no storage requirements in excess of just the raw text in the files.

Data Processing

To transform the data into the format previously described, some data processing is required. All programs that we used to transform the datasets are available at <https://github.com/mbd-doc2vec-team/mbd-doc2vec/tree/master/data-processing>.

IMDB Dataset

The IMDB dataset needed minimal processing. The bulk of the work had already been performed by the authors of the dataset. Beyond what had already been done, punctuation, symbols and elements of markup were removed, and each document concatenated into a file as discussed in the previous section.

Wikipedia Dataset

The Wikipedia dataset required extensive pre-processing, with each document needing to be parsed from the XML, and the plain text extracted from the wiki markup. Additionally, "fake" documents (e.g. portals, category pages) were removed. Each wiki page was then tokenized and stored in a file as previously stated.

Chromium Dataset

The Chromium codebase required extensive pre-processing as well, with each C++ file being tokenized and having unicode characters (e.g. `̀`) removed. Unlike the other two datasets, because punctuation and uppercase/lowercase distinctions are semantically significant in code, we retained punctuation and the original casing in tokens. Each comment and string literal was then tokenized as well, with punctuation being removed from these tokens only, and the resulting tokens being stored as discussed in the previous section. For an example of this process, the document

```
// Frobulates the bar, in an efficient manner.  
void Frobulate(Bar &bar);
```

would be represented by a line in the file as follows

```
Frobulates the bar in an efficient manner void Frobulate ( Bar & bar ) ;
```

Experiments

Following are details as to how we carried out our experiments. Full source code to do this can be found at <https://github.com/mbd-doc2vec-team/mbd-doc2vec>, and full datasets can be obtained as listed in the “Data Sources” section of this paper.

IMDB Dataset

The experimentation was undertaken with the method as follows.

Preprocessing

- The entire set of 100,000 documents were processed as described in the previous sections. The reason for concatenating all documents into a single file is that the dataset is small and can comfortably fit within memory, so it does not suffer from this simplification.

Training

- One of the co-authors of the original paper, Mikolov, provided some guidance and recommended the use of gensim’s skip-gram model, which is trained for 20 iterations on the entire corpus of reviews.
- After each iteration, the order of the documents is randomized.

Evaluation

- The training document vectors are paired with their labels and a logistic regression is calculated to fit the training data.
- Prediction of sentiment using the testing document vectors takes place and the accuracy of the prediction versus the actual labels is reported.

Wikipedia Dataset

Chromium Dataset

Results

IMDB Dataset

Through experimentation and switching from a neural network based approach to a logistic regression, the final best accuracy achieved whilst attempting to replicate the results was 89.4%. Being unable to replicate the results of Le and Mikolov in [3], some research into the result was conducted. Many others were trying their hardest to replicate these results, but ultimately the co-author Mikolov, who was not involved in the experimentation found the issue.

This was reported in [13], where Mikolov states the 92.6% accuracy is achievable using hierarchical softmax “only when the training and test data are not shuffled. Thus, we consider this result to be invalid”. We therefore confirm this non-replication.

Wikipedia Dataset

Experimentation with the Wikipedia dataset reveals some remarkable semantic connections. Table 1 shows articles with semantic contexts that most closely resemble the article describing the *XKCD* webcomic. For example, Randal Munroe is the author of the *XKCD* webcomic, while *MS Paint Adventures*, *Penny Arcade* and *Ctrl+Alt+Del* are all webcomics too.

Page name	Cosine similarity
randall munroe	0.517820239067
ms paint adventures	0.443829506636
time (xkcd)	0.440876543522
penny arcade	0.424748897552
alexey pajitnov	0.424251556396
ctrl+alt+del (webcomic)	0.423852056265

Table 1: Wikipedia pages most similar to *XKCD*

A feature of representing the embedding in a vector, is being able to find A in the relationship “X is to Y, as A is to B” by summing the vectors for X and B, and subtracting Y. A simple example of how this works is as follows for “Bigger is to Big, as _____ is to Large”

```
>>> model.most_similar([model['large'] + model['bigger'] - model['big']])
[(u'larger', 0.7768259048461914)]
```

Citizen Kane is a film widely and highly regarded by critics, and as such it has become a part of popular culture to ask the question: “What compares to the critical acclaim of *Citizen Kane* in other types of entertainment?”. To explore the Wikipedia model’s embedding of semantics, the model was applied to answer the question, “What is the *Citizen Kane* of video games?”, a question widely debated on the internet.

```
>>> model.docvecs.most_similar([citizen_kane, video_game], [film])
[('chrono_trigger', 0.5553357601165771),
 ('day_of_the_tentacle', 0.5282092094421387),
 ('deus_ex', 0.5077250599861145),
 ('kid_icarus', 0.506332516670227),
 ('sid_meier's_alpha_centauri', 0.5028576850891113),
 ('cyberspace', 0.49190062284469604),
```

```
(‘chrono_cross’, 0.4834246337413788),
(‘industrial_and_organizational_psychology’, 0.4763486087322235),
(‘diablo_ii’, 0.4696195721626282),
(‘content-control_software’, 0.46697914600372314)]
```

‘Cyberspace’, ‘Industrial and Organizational Psychology’ and ‘Content-Control Software’ are examples of results that have semantic similarities between *Citizen Kane* and video games but are not video games titles. This method is not a flawless way of generating human-comprehensible relationships.

Chromium Dataset

The final application of this document vector model was to investigate the semantic relationships it gathers when applied to source code files. The results were mixed. Examples of clusters included platform-specific implementations and their header files, memory allocation headers from the likes of Webkit and SQLite, GUI components from Webkit and third-party libraries and various message dispatchers. Otherwise, there aren’t a lot of relations that can be seen in the overall distribution.

Conclusion

Future Work

References

- [1] Yin Zhang, Rong Jin, and Zhi-Hua Zhou. “Understanding bag-of-words model: a statistical framework”. In: *International Journal of Machine Learning and Cybernetics* 1.1-4 (2010), pp. 43–52.
- [2] Tomas Mikolov et al. “Efficient estimation of word representations in vector space”. In: *arXiv preprint arXiv:1301.3781* (2013).
- [3] Quoc V Le and Tomas Mikolov. “Distributed Representations of Sentences and Documents.” In: *ICML*. Vol. 14. 2014, pp. 1188–1196.
- [4] *Google Groups*. URL: <https://groups.google.com/d/msg/word2vec-toolkit/q49firqnoqro/bp--14e4unwj>.
- [5] David E Rumelhart, Geoffrey E Hinton, and Ronald J Williams. “Learning representations by back-propagating errors”. In: *Cognitive modeling* 5.3 (1988), p. 1.
- [6] Yoshua Bengio et al. “A neural probabilistic language model”. In: *journal of machine learning research* 3.Feb (2003), pp. 1137–1155.
- [7] Tomas Mikolov et al. “Distributed representations of words and phrases and their compositionality”. In: *Advances in neural information processing systems*. 2013, pp. 3111–3119.
- [8] Jey Han Lau and Timothy Baldwin. “An Empirical Evaluation of doc2vec with Practical Insights into Document Embedding Generation”. In: *arXiv preprint arXiv:1607.05368* (2016).
- [9] Andrew M Dai, Christopher Olah, and Quoc V Le. “Document embedding with paragraph vectors”. In: *arXiv preprint arXiv:1507.07998* (2015).
- [10] Andrew L. Maas et al. “Learning Word Vectors for Sentiment Analysis”. In: *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*. Portland, Oregon, USA: Association for Computational Linguistics, 2011, pp. 142–150. URL: <http://www.aclweb.org/anthology/P11-1015>.

- [11] Meta. *Data dump torrents — Meta, discussion about Wikimedia projects*. [Online; accessed 9-August-2016]. 2016. URL: https://meta.wikimedia.org/w/index.php?title=Data_dump_torrents.
- [12] *Git repositories on chromium*. URL: <https://chromium.googlesource.com/>.
- [13] Grégoire Mesnil et al. “Ensemble of generative and discriminative techniques for sentiment analysis of movie reviews”. In: *arXiv preprint arXiv:1412.5335* (2014).